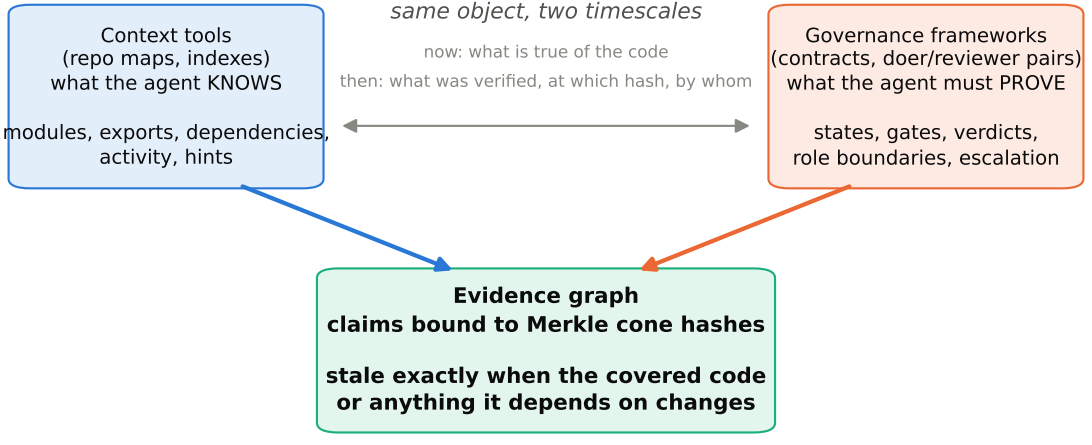
Context tools
(repo maps, indexes)
what the agent KNOWS
modules, exports, dependencies,
activity, hints
same object, two timescales
now: what is true of the code
then: what was verified, at which hash, by whom
Governance frameworks
(contracts, doer/reviewer pairs)
what the agent must PROVE
states, gates, verdicts,
role boundaries, escalation
Evidence graph
claims bound to Merkle cone hashes
stale exactly when the covered code
or anything it depends on changes
outputs: situation brief (tokens), signed ledger (audit), mechanical gate (merge)